



**University of Dhaka**  
**Department of Robotics and Mechatronics Engineering**

Course Code: RME 4100 (Project)

**Adaptive Neural Control for Mars Rover Wheel  
Entrapment Recovery in Granular Terrain**

Meraj Hossain Promit[SH-172-029]

&

Chandak Chakma[JN-172-026]

**Under the supervision of**

---

**Supervisor**

Md. Jubair Ahmed Surov  
Lecturer

Department of Robotics and Mechatronics  
Engineering  
University of Dhaka

---

**Co-Supervisor**

Dr. Sejuti Rahman  
Associate Professor

Department of Robotics and Mechatronics  
Engineering  
University of Dhaka

### Abstract

Wheel entrapment in granular regolith is one of the most critical failure modes for planetary rovers, as demonstrated by the permanent immobilization of NASA’s Spirit rover in 2009 [14]. This project develops an autonomous entrapment recovery system for a 6-wheel rocker-bogie Mars rover using deep reinforcement learning trained entirely in simulation.

We couple the Newton physics engine’s Material Point Method (MPM) [21] with Isaac Lab’s [9] rigid-body solver to produce physically accurate granular soil behaviour under Martian gravity ( $3.72 \text{ m/s}^2$ ). A recurrent policy trained with Proximal Policy Optimization (PPO) [2] and Gated Recurrent Units (GRU) [13] operates on a 29-dimensional proprioceptive state comprising wheel velocities, slip ratios, IMU readings, drive torques, and dual entrapment detection flags. Curriculum learning progressively increases wheel sinkage from 12 cm to 22 cm across 32 parallel training environments on a single workstation (Intel i7-1255U, NVIDIA RTX 5070 Ti).

Training over 200,000 timesteps achieves an 60% escape rate with a mean forward velocity of 0.25 m/s and full curriculum completion. The policy develops emergent escape behaviours including coordinated rocking manoeuvres and asymmetric wheel-steering sequences not explicitly programmed. Planned next steps cover scaling to 2M+ timesteps, a CNN-GRU sinkage detection classifier, and on-rover deployment via sim2sim validation or possibly ONNX Runtime on a Raspberry Pi 5 at 10 Hz.

# Contents

|          |                                                                                |           |
|----------|--------------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                                            | <b>6</b>  |
| 1.1      | Problem Statement . . . . .                                                    | 6         |
| 1.2      | Problem Definition . . . . .                                                   | 7         |
| 1.3      | Motivation: The Spirit Rover Incident (2009) . . . . .                         | 8         |
| 1.4      | Objectives . . . . .                                                           | 9         |
| <b>2</b> | <b>Related Work</b>                                                            | <b>10</b> |
| 2.1      | Foundational Concepts and Terminology . . . . .                                | 10        |
| 2.2      | Escape Strategies for Wheeled Robots . . . . .                                 | 11        |
| 2.2.1    | Bi and Ding (2026): DRL-Based Escape for Skid-Steered Mobile Robots . . . . .  | 11        |
| 2.2.2    | Inotsume et al. (2020): Slip Prediction for Planetary Rovers . . . . .         | 13        |
| 2.2.3    | Ding et al. (2011): Wheel-Terrain Interaction Modeling . . . . .               | 13        |
| 2.2.4    | Jain et al. (2019): RL for Planetary Rover Locomotion . . . . .                | 14        |
| 2.2.5    | Fankhauser et al. (2018): ANYmal Recovery from Falls . . . . .                 | 14        |
| 2.2.6    | Mortensen and Bøgh (2024): RL RoverLab RL Suite for Planetary Rovers . . . . . | 14        |
| 2.3      | Granular Media Simulation . . . . .                                            | 15        |
| 2.4      | Summary and Comparison . . . . .                                               | 15        |
| 2.5      | Gaps in Existing Work . . . . .                                                | 16        |
| <b>3</b> | <b>Methodology</b>                                                             | <b>18</b> |
| 3.1      | Work Already Done . . . . .                                                    | 22        |
| 3.1.1    | Simulation Environment . . . . .                                               | 22        |
| 3.1.2    | Robot Platform . . . . .                                                       | 24        |
| 3.1.3    | Observation and Action Spaces . . . . .                                        | 25        |
| 3.1.4    | Dual-Sensor Entrapment Detection . . . . .                                     | 25        |
| 3.1.5    | Reward Function . . . . .                                                      | 26        |
| 3.1.6    | PPO Training Configuration . . . . .                                           | 27        |
| 3.1.7    | Training Hardware . . . . .                                                    | 27        |
| 3.1.8    | Curriculum Learning . . . . .                                                  | 28        |
| 3.1.9    | Reward Tuning and Bug Fixes . . . . .                                          | 28        |
| 3.2      | Expected Work to Be Done . . . . .                                             | 29        |
| 3.2.1    | Phase 1: CNN-GRU Sinkage Detection Classifier . . . . .                        | 29        |
| 3.2.2    | Scaling Training . . . . .                                                     | 29        |
| 3.2.3    | Phase 3: Sim-to-Real Deployment . . . . .                                      | 29        |
| 3.3      | Full System Architecture: Dual-Brain Policy . . . . .                          | 30        |
| <b>4</b> | <b>Results</b>                                                                 | <b>31</b> |
| 4.1      | Training Overview . . . . .                                                    | 31        |
| 4.2      | Key Performance Metrics . . . . .                                              | 32        |
| 4.3      | Reward Analysis . . . . .                                                      | 32        |

|          |                                                                   |           |
|----------|-------------------------------------------------------------------|-----------|
| 4.4      | Environment Metrics . . . . .                                     | 33        |
| 4.5      | PPO Training Losses . . . . .                                     | 33        |
| 4.6      | Episode Running Data . . . . .                                    | 33        |
| 4.7      | Entrapment Anomaly Visualization . . . . .                        | 34        |
| 4.8      | Observed Escape Behaviors . . . . .                               | 35        |
| 4.9      | Comparison with Bi and Ding (2026) . . . . .                      | 35        |
| <b>5</b> | <b>Conclusion</b>                                                 | <b>36</b> |
| 5.1      | Summary of Contributions . . . . .                                | 36        |
| 5.2      | Preliminary Results . . . . .                                     | 36        |
| 5.3      | Future Work . . . . .                                             | 36        |
| 5.3.1    | Priority 1: Immediate Next Steps . . . . .                        | 36        |
| 5.3.2    | Priority 2: Sim-to-Sim Validation . . . . .                       | 37        |
| 5.3.3    | Priority 3: Sim-to-Real Transfer ( <i>if possible</i> ) . . . . . | 37        |

## List of Figures

|    |                                                                                                                                                                                                                                                                                                                                                                                                             |    |
|----|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 1  | Timeline of the Spirit Rover entrapment incident. After 5 years of successful exploration, Spirit became permanently immobilized in soft regolith, motivating autonomous entrapment recovery research. . . . .                                                                                                                                                                                              | 8  |
| 2  | Escape policy generation method. Top: the training pipeline in simulation with reward shaping and curriculum learning. Bottom: the deployment pipeline on the physical or simulated Mars rover. . . . .                                                                                                                                                                                                     | 18 |
| 3  | Simulation framework. The framework couples 6-wheel MPM particle interactions with granular sinkage via Newton MPM, and incorporates curriculum-driven sinkage depth randomization. . . . .                                                                                                                                                                                                                 | 19 |
| 4  | Structure of the Mars rover control policy network (GRU, 256 hidden units, 29D input, 10D output). . . . .                                                                                                                                                                                                                                                                                                  | 20 |
| 5  | Block diagram of the reward-shaped DRL training loop. The curriculum scheduler updates environment difficulty and reward weights. The closed-loop interaction between PPO and the environment generates transitions, which are stored and sampled as minibatches for policy updates. . . . .                                                                                                                | 21 |
| 6  | Physics simulation stack. The key innovation is two-way coupling between rigid-body dynamics and granular MPM particles. . . . .                                                                                                                                                                                                                                                                            | 22 |
| 7  | Newton simulation environment showing the AAU 6-wheel Mars rover on granular regolith. The MPM particle bed ( $2.0\text{ m} \times 2.0\text{ m} \times 0.15\text{ m}$ ) provides realistic sinkage and wheel-sand interaction under Martian gravity ( $3.72\text{ m/s}^2$ ). Screenshots to be captured from Newton ViewerGL. . . . .                                                                       | 22 |
| 8  | The AAU Mars rover platform used in simulation. (a) USD model render showing the 6-wheel rocker-bogie suspension from RLRoverLab. (b) Rover in Newton simulation environment with granular regolith particles under Martian gravity ( $3.72\text{ m/s}^2$ ). . . . .                                                                                                                                        | 24 |
| 9  | Planned CNN-GRU sinkage detection pipeline. . . . .                                                                                                                                                                                                                                                                                                                                                         | 29 |
| 10 | Planned sim-to-real deployment pipeline. . . . .                                                                                                                                                                                                                                                                                                                                                            | 29 |
| 11 | Dual-brain onboard architecture deployed on the Raspberry Pi 5. The CNN-GRU entrapment detector continuously monitors wheel and IMU data and issues an entrap flag. During normal traversal the <b>Navigation Brain</b> (RLRoverLab policy) drives the rover; when entrapment is detected the <b>Escape Brain</b> (our GRU-PPO policy) takes over until the rover escapes, then hands control back. . . . . | 30 |
| 12 | Training overview across 200K timesteps: reward convergence, escape rate, forward velocity, slip ratio, losses, and exploration metrics. GRU-PPO (red) vs. feedforward baseline (blue). . . . .                                                                                                                                                                                                             | 31 |
| 13 | Reward breakdown showing min, mean, and max episode rewards over training. . . . .                                                                                                                                                                                                                                                                                                                          | 32 |
| 14 | Environment metrics: forward velocity, slip ratio, entrapment rate, and torque anomaly rate. . . . .                                                                                                                                                                                                                                                                                                        | 33 |
| 15 | PPO training losses: policy loss, value loss, and entropy loss. . . . .                                                                                                                                                                                                                                                                                                                                     | 33 |

|    |                                                                                                                                                                                                                                         |    |
|----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 16 | Episode running data dashboard for a single evaluation episode with the trained GRU-PPO policy (to be generated). Format matches Bi and Ding [1], Figs. 18/19.                                                                          | 34 |
| 17 | The Mars rover entrapment anomaly in the Newton simulation (analogous to Bi and Ding [1], Fig. 13). (a) Rover wheels sunk into granular regolith bed. (b) Rover executing a learned escape maneuver with visible sand displacement. . . | 34 |

## List of Tables

|    |                                                                                                                                                                                                                                                                                                                                                                                                                            |    |
|----|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----|
| 1  | Reward function v4 used during PPO training. The escape bonus (weight 20) dominates during entrapment, while the rocking bonus encourages active self-extraction. Compared to Bi & Ding (2026) [1], who use GAIL-fused [31] rewards ( $R = 0.9R_H + 0.1R_D$ ) with expert demonstrations, our reward is fully handcrafted and designed for 6-DOF steering under Martian gravity without requiring any expert data. . . . . | 7  |
| 2  | Comparison of related works on wheeled robot entrapment and recovery . . .                                                                                                                                                                                                                                                                                                                                                 | 16 |
| 3  | Environment configuration parameters . . . . .                                                                                                                                                                                                                                                                                                                                                                             | 23 |
| 4  | AAU Mars rover joint configuration . . . . .                                                                                                                                                                                                                                                                                                                                                                               | 24 |
| 5  | Observation space components (29D total) . . . . .                                                                                                                                                                                                                                                                                                                                                                         | 25 |
| 6  | Reward function components (v3 methodology reference) . . . . .                                                                                                                                                                                                                                                                                                                                                            | 26 |
| 7  | PPO training hyperparameters and GRU architecture . . . . .                                                                                                                                                                                                                                                                                                                                                                | 27 |
| 8  | Training system configuration (Lenovo ThinkBook 14 G4 IAP) . . . . .                                                                                                                                                                                                                                                                                                                                                       | 28 |
| 9  | Training results at 200,000 timesteps (RTX 5070 Ti, 64 parallel environments) . .                                                                                                                                                                                                                                                                                                                                          | 32 |
| 10 | Comparison with Bi and Ding (2026) [1] . . . . .                                                                                                                                                                                                                                                                                                                                                                           | 35 |

# 1 Introduction

On May 1, 2009, NASA’s Spirit rover drove into soft sulfate-rich soil and never moved again. Eight months of recovery attempts, thousands of ground-based tests, and 150 million kilometres of communication delay could not save it. This project asks: **what if the rover could have saved itself?**

## 1.1 Problem Statement

Planetary exploration rovers operating on Mars and other celestial bodies [15, 16, 17] face a critical operational challenge: **wheel entrapment in granular regolith**. When a rover’s wheels sink into loose sand or soil, the vehicle loses mobility and may become permanently immobilized. This problem is particularly acute for several reasons:

- **Rocker-bogie suspension:** While excellent for rough terrain, it distributes weight in ways that can cause individual wheels to sink disproportionately in soft soil.
- **Sandy and cratered terrain:** The Martian surface contains extensive dune fields, impact craters filled with fine regolith, and wind-blown sand deposits that present persistent mobility hazards.
- **Long communication delays:** With Earth–Mars one-way latencies of 4–24 minutes, real-time teleoperation for recovery is impractical, necessitating fully autonomous solutions.
- **Extended mission lifetimes:** As missions extend beyond planned durations, rovers inevitably encounter terrain types not anticipated during original mission planning.

Wheel entrapment is formally characterised by the simultaneous satisfaction of three conditions:

$$f_{\text{ent}} = 1 \iff \begin{cases} |v_x| < 0.15 \text{ m/s} & \text{(near-zero forward velocity)} \\ \bar{\sigma}_{\text{slip}} > 0.4 & \text{(high mean wheel slip)} \\ d_{\text{sinkage}} > 0 & \text{(positive wheel sinkage)} \end{cases} \quad (1)$$

where  $v_x$  is the body-frame forward velocity,  $\bar{\sigma}_{\text{slip}} = \frac{1}{6} \sum_i \sigma_i$  is the mean slip ratio across all six drive wheels, and  $d_{\text{sinkage}}$  is the wheel sinkage depth into the granular medium. All three conditions must hold for at least **15 consecutive steps** (0.6 s) before the entrapment flag  $f_{\text{ent}}$  is raised.

## 1.2 Problem Definition

The entrapment recovery problem is formulated as a **partially observable Markov decision process (POMDP)**. At each timestep  $t$ , the rover receives a proprioceptive observation  $\mathbf{o}_t \in \mathbb{R}^{29}$  and must select an action  $\mathbf{a}_t \in \mathbb{R}^{10}$  to maximise the expected discounted return:

$$J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \gamma^t R(\mathbf{s}_t, \mathbf{a}_t) \right], \quad \gamma = 0.99 \quad (2)$$

The shaped reward  $R$  is the sum of seven terms : three incentives and four penalties, designed to reward escape while discouraging dangerous manoeuvres:

$$R_t = \underbrace{r_{\text{progress}} + r_{\text{escape}} + r_{\text{rocking}}}_{\text{incentives}} - \underbrace{(p_{\text{slip}} + p_{\text{tilt}} + p_{\text{smooth}} + p_{\text{abnormal}})}_{\text{penalties (all positive, subtracted)}} \quad (3)$$

| Term                  | Weight | Formula & Description                                                                                                                                                     |
|-----------------------|--------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Incentives</b>     |        |                                                                                                                                                                           |
| $r_{\text{progress}}$ | +1.5   | $v_x \cdot \Delta t \cdot (1 - f_{\text{ent}})$ (forward velocity reward, suppressed when trapped so rocking takes over)                                                  |
| $r_{\text{escape}}$   | +20.0  | Sparse milestones at 0.5 m ( $\times 0.2$ ), 1.0 m ( $\times 0.4$ ), 1.5 m ( $\times 1.0$ ) scaled by $\tau_{\text{time}}$ ; plus dense shaping $0.5 d \Delta t$ per step |
| $r_{\text{rocking}}$  | +2.0   | $\Delta v_x \cdot f_{\text{ent}} \cdot \Delta t$ (rewards velocity change during entrapment to encourage active rocking manoeuvres)                                       |
| <b>Penalties</b>      |        |                                                                                                                                                                           |
| $p_{\text{slip}}$     | 0.5    | $\bar{\sigma}_{\text{slip}} \cdot (1 - f_{\text{ent}}) \cdot \Delta t$ (penalises excessive wheel slip during normal driving)                                             |
| $p_{\text{tilt}}$     | 0.3    | $\ \boldsymbol{\omega}_{xy}\ _2 \cdot \Delta t$ (penalises roll/pitch angular velocity to maintain rover stability)                                                       |
| $p_{\text{smooth}}$   | 0.05   | $\ \mathbf{a}_t - \mathbf{a}_{t-1}\ _2 \cdot \Delta t$ (penalises jerky action changes to encourage smooth control)                                                       |
| $p_{\text{abnormal}}$ | 0.3    | $(-v_x)^+ \cdot f_{\text{anomaly}} \cdot (1 - f_{\text{ent}}) \cdot \Delta t$ (penalises backward motion under torque anomaly when not trapped)                           |

$f_{\text{ent}}$ : entrapment flag;  $f_{\text{anomaly}}$ : torque anomaly flag;  $d$ : distance from spawn;  $\Delta t = 0.04$  s

Table 1: Reward function v4 used during PPO training. The escape bonus (weight 20) dominates during entrapment, while the rocking bonus encourages active self-extraction. Compared to Bi & Ding (2026) [1], who use GAIL-fused [31] rewards ( $R = 0.9R_H + 0.1R_D$ ) with expert demonstrations, our reward is fully handcrafted and designed for 6-DOF steering under Martian gravity without requiring any expert data.

The partial observability arises because the rover cannot directly sense the depth of sink-



age per wheel, the local terrain composition, or its position relative to the entrapment zone boundary. This motivates a **recurrent policy** based on the Gated Recurrent Unit (GRU) [13] that maintains a hidden state  $\mathbf{h}_t$  to implicitly estimate these unobserved variables from the history of observations. GRUs are a lightweight alternative to LSTMs [32] that achieve comparable performance on temporal sequence modeling with fewer parameters.

### 1.3 Motivation: The Spirit Rover Incident (2009)

The most prominent example of wheel entrapment in planetary exploration occurred on May 1, 2009, when NASA’s Mars Exploration Rover *Spirit* became embedded in soft sulfate-rich regolith at a site named “Troy” near the western edge of Home Plate in Gusev Crater [14, 15]. Figure 1 illustrates the timeline of this critical incident.




Figure 1: Timeline of the Spirit Rover entrapment incident. After 5 years of successful exploration, Spirit became permanently immobilized in soft regolith, motivating autonomous entrapment recovery research.

Key lessons from the Spirit incident that inform this research:

1. **Detection was delayed:** Spirit had driven partway into the soft soil before operators recognized the hazard, as slip ratio was not monitored in real-time.
2. **Recovery required months:** The 8-month extraction campaign involved extensive Earth-based testing on rover replicas, constrained by communication delays.
3. **Autonomous response was absent:** Spirit had no onboard capability to detect or respond to entrapment autonomously.
4. **Granular physics are critical:** The behavior of the sulfate-rich soil was poorly understood, leading to extraction commands that inadvertently worsened the entrapment.

## 1.4 Objectives

This research aims to develop a complete autonomous entrapment detection and recovery system for Mars rovers. The five specific objectives are:

**1. Realistic Simulation:** Build a high-fidelity simulation environment coupling granular material physics (Newton MPM) with articulated rover dynamics (Isaac Lab) under Martian gravity ( $3.72 \text{ m/s}^2$ ), enabling realistic wheel–soil interaction for RL training.

**2. Entrapment Detection:** Develop a dual-sensor detection scheme combining slip-ratio monitoring and drive-torque anomaly detection for robust, real-time entrapment classification without external sensing or GPS.

**3. Recovery Policy:** Train a recurrent RL policy (PPO + GRU) that autonomously discovers and executes effective escape manoeuvres, including rocking, steering pivots, and differential drive, from varying degrees of wheel sinkage (6–12 cm curriculum).

**4. Curriculum Learning:** Implement progressive difficulty scheduling that begins with shallow sinkage and gradually increases entrapment severity, enabling the policy to learn robust escape strategies that generalise across terrain conditions.

**5. Sim-to-Real Transfer:** Establish a complete deployment pipeline comprising domain randomisation during training, ONNX export of the trained GRU policy, and real-time inference on a Raspberry Pi 5, bridging the simulation-to-reality gap for onboard autonomous operation.

## 2 Related Work

This section reviews prior work in four key areas: (1) foundational concepts and terminology, (2) wheeled robot entrapment and escape strategies, (3) granular media simulation for robotics, and (4) reinforcement learning for locomotion and recovery. We conclude with a summary table and an explicit identification of the gaps our work addresses.

### 2.1 Foundational Concepts and Terminology

This subsection introduces the core technical concepts underpinning our system for readers less familiar with one or more of the constituent domains.

#### Rocker-Bogie Suspension

The rocker-bogie suspension is a passive linkage mechanism used on all NASA Mars rovers (Sojourner, Spirit, Opportunity, Curiosity, Perseverance) and our AAU platform. It consists of a *rocker* arm connecting the front wheel to a *bogie* sub-assembly bearing the middle and rear wheels. A differential link couples the two rocker arms so that the body roll is kept to half the terrain-induced tilt of any single rocker. This passive compliance allows the six wheels to maintain ground contact over obstacles up to twice the wheel diameter. The kinematic complexity, arising from passive joints, a differential, and no central actuator for body attitude, makes control significantly harder than a simple 4-wheel skid-steered platform.

#### Wheel–Soil Interaction and Slip Ratio

When a driven wheel rotates in granular soil, the soil deforms and the wheel may sink (*sinkage*) while the contact patch shears the granular medium. The *slip ratio*  $s \in [0, 1]$  quantifies this mismatch:

$$s = 1 - \frac{v_{\text{linear}}}{\omega r} \quad (4)$$

where  $v_{\text{linear}}$  is the measured forward velocity,  $\omega$  the wheel angular velocity, and  $r$  the wheel radius. A slip ratio approaching 1 indicates the wheel is spinning freely with no forward progress, which is a hallmark of entrapment. Sustained high slip combined with low forward velocity is the primary detection signal used in this work.

#### Material Point Method (MPM)

MPM [21] is a hybrid Eulerian–Lagrangian continuum method that represents the material as a set of *Lagrangian particles* (tracking mass, momentum, and deformation gradient) and a background *Eulerian grid* (for solving the equations of motion). At each timestep, particle data is transferred to the grid (*P2G*), grid equations are solved, and updated velocities are mapped back to particles (*G2P*). For granular sand, the Drucker–Prager yield surface [23] governs plastic flow: material flows when the stress state reaches the yield cone, and elastically recovers otherwise. Compared to rigid-body approximations, MPM correctly models sinkage

progression, bulldozing resistance, and particle redistribution during wheel spinning, all of which are critical for entrapment simulation.

### Proximal Policy Optimization (PPO)

PPO [2] is an on-policy actor-critic RL algorithm that improves training stability by constraining the policy update step size via a clipped surrogate objective:

$$\mathcal{L}^{\text{CLIP}}(\theta) = \mathbb{E}_t [\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t)] \quad (5)$$

where  $r_t(\theta) = \pi_\theta(a_t|s_t)/\pi_{\theta_{\text{old}}}(a_t|s_t)$  is the probability ratio and  $\hat{A}_t$  is the generalized advantage estimate [27]. The clip parameter  $\epsilon$  (typically 0.2) prevents destructively large updates. PPO has become the default algorithm for continuous-control locomotion tasks due to its simplicity, sample efficiency, and robustness to hyperparameter choice.

### Gated Recurrent Unit (GRU)

A GRU [13] is a recurrent neural network cell that maintains a hidden state  $\mathbf{h}_t$  summarising history:

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t \quad (6)$$

where  $\mathbf{z}_t$  is the update gate and  $\tilde{\mathbf{h}}_t$  is the candidate hidden state. GRUs are preferred over LSTMs [32] when model size is constrained (e.g., deployment on a Raspberry Pi 5), as they achieve comparable sequence-modelling performance with roughly one-third fewer parameters. In our setting, the GRU enables the policy to track slow-developing entrapment dynamics (sinkage over several seconds) that would be invisible to a memoryless MLP policy.

## 2.2 Escape Strategies for Wheeled Robots

### 2.2.1 Bi and Ding (2026): DRL-Based Escape for Skid-Steered Mobile Robots

Bi and Ding [1] presented a deep reinforcement learning framework for escape and path point tracking of skid-steered mobile robots (SSMRs) under undulating terrain conditions in outdoor orchards. Their work is the most directly relevant to ours.

**Description.** The authors developed a two-phase system: (1) an abnormal state detection module using motor torque analysis, and (2) a DRL-based escape policy using PPO with LSTM networks [32]. Their key innovation is a *reward fusion* strategy combining handcrafted rewards with GAIL-derived [31] (Generative Adversarial Imitation Learning) discriminator rewards:  $R = 0.9R_H + 0.1R_D$ , where  $R_H$  includes path tracking, abnormal action, and conditional rewards, and  $R_D$  comes from a discriminator trained on expert demonstrations. The system was trained for 4 million timesteps in a Unity simulation environment with extensive domain randomization (20+ parameters, see their Table 9) and validated in real-world outdoor orchard conditions, achieving 90% escape success rate in both simulation and field tests.

Their observation space comprised 30 dimensions organized as  $\mathbf{s} = \{\mathbf{s}_{\text{motor}}, \mathbf{s}_{\text{motion}}, \mathbf{s}_{\text{position}}, \mathbf{s}_{\text{abnormal}}\}$ : 16D motor data (set speed, current speed, current torque, average torque  $\times 4$  motors), 7D motion state (quaternion attitude + angular velocity), 2D position (distance and angle to target), and 1D anomaly flag. The action space is 2D: left and right wheel velocities  $\mathbf{a} = [\omega_l, \omega_r]$ .

### Strengths.

- Achieved 90% escape success rate in both simulation and real-world tests, demonstrating strong sim-to-real transfer via domain randomization.
- LSTM-based temporal modeling captures the sequential nature of entrapment dynamics, enabling memory-dependent escape strategies.
- GAIL reward fusion ( $R = 0.9R_H + 0.1R_D$ ) combines human intuition with learned reward shaping, producing smoother and more natural escape trajectories compared to pure handcrafted or pure imitation approaches.
- Comprehensive anomaly detection based on motor torque threshold (Eq. 10 in their paper):  $\frac{1}{c_t} \sum_{n=t-c_t}^t \frac{|T_n|}{T_N} > T_{\text{thread}}$  with  $T_{\text{thread}} = 0.95$ .
- Validated on real outdoor terrain (orchards) with undulating surfaces using UWB positioning.

### Weaknesses.

- Limited to 4-wheel skid-steered platforms; does not address the more complex kinematics of 6-wheel rocker-bogie suspension systems used in Mars rovers (Curiosity, Perseverance).
- Uses rigid terrain simulation (Unity) and does not model granular soil mechanics (particle flow, sinkage dynamics, bulldozing resistance) that are critical for planetary regolith interaction.
- Tested under Earth gravity ( $9.81 \text{ m/s}^2$ ), not Mars gravity ( $3.72 \text{ m/s}^2$ ), which significantly affects soil-wheel traction and sinkage behavior.
- Requires expert demonstrations for GAIL training, which may not be available for novel Mars terrain conditions.
- Two-dimensional action space ( $\omega_l, \omega_r$ ) does not leverage independent steering joints for more sophisticated escape maneuvers.
- Entrapment is caused by tire lift-off on undulating terrain, not by granular sinkage, which is a fundamentally different failure mode.

### 2.2.2 Inotsume et al. (2020): Slip Prediction for Planetary Rovers

**Description.** Inotsume et al. [3] investigated real-time slip prediction for planetary rovers using proprioceptive sensing and terrain parameter estimation. A Kalman filter fuses motor current, wheel sinkage estimates, and terrain slope to predict slip ratio based on Bekker–Wong terramechanics theory.

#### Strengths.

- Principled physics-based approach grounded in terramechanics theory; provides interpretable terrain parameter estimates.
- Runs in real-time on embedded hardware with low computational requirements.
- Directly applicable to Mars rover missions with existing sensor suites.

#### Weaknesses.

- Focuses on slip prediction only; does not address active recovery from entrapment.
- Requires careful calibration for each terrain type, limiting generalization to unknown Martian soils.
- Simplified Bekker–Wong model may not capture complex particle–wheel coupling during deep sinkage.

### 2.2.3 Ding et al. (2011): Wheel–Terrain Interaction Modeling

**Description.** Ding et al. [4] presented a comprehensive wheel–soil interaction model for planetary rovers, deriving analytical expressions for drawbar pull, rolling resistance, and sinkage as functions of wheel geometry, soil parameters (cohesion, internal friction angle, shear deformation modulus), and slip ratio. The model accounts for multi-pass effects where trailing wheels encounter pre-compacted soil.

#### Strengths.

- Rigorous analytical framework validated experimentally with single-wheel testbed data.
- Accounts for multi-wheel interactions and soil compaction by preceding wheels.

#### Weaknesses.

- Steady-state analysis only; does not capture transient dynamics during entrapment onset or recovery.
- Assumes homogeneous soil properties, unrealistic for heterogeneous Martian terrain.
- Does not address control strategies for escape from entrapment.

### 2.2.4 Jain et al. (2019): RL for Planetary Rover Locomotion

**Description.** Jain et al. [5] applied DQN for adaptive locomotion of planetary rovers, learning to select discrete gaits (forward, turn, reverse) based on terrain features from onboard cameras in Gazebo simulation.

**Strengths.** Demonstrated that RL can learn terrain-adaptive locomotion; uses visual observations relevant to missions.

**Weaknesses.** Discrete action space limits expressiveness of recovery maneuvers; does not specifically address entrapment; approximate terrain deformation model.

### 2.2.5 Fankhauser et al. (2018): ANYmal Recovery from Falls

**Description.** Fankhauser et al. [6] developed a recovery controller for the ANYmal quadruped robot combining optimization-based motion planning with learned recovery policies. The multi-phase architecture (detect → recover → resume) provides a useful design template.

**Strengths.** Successful sim-to-real transfer for dynamic recovery behaviors; multi-phase architecture is applicable to our problem.

**Weaknesses.** Designed for legged robots with body reconfiguration; wheel entrapment in granular media presents fundamentally different physical constraints.

### 2.2.6 Mortensen and Bøgh (2024): RL RoverLab RL Suite for Planetary Rovers

Mortensen and Bøgh [7] introduced RL RoverLab, an open-source reinforcement learning suite for simulating and training planetary rovers on extraterrestrial bodies. This work is directly relevant because our project builds upon the AAU Mars Rover assets, training frameworks, and pre-trained navigation policies provided by RL RoverLab.

**Description.** RL RoverLab is implemented on top of the Nvidia ORBIT framework [39], which uses Isaac Sim as the physics and graphics engine. The suite provides: (1) four terrain types (flat, hill, obstacle, maze) with procedural generation; (2) three rover models: the AAU Rover Complex (6-wheel, rocker-bogie, ~1 m footprint), AAU Rover Simple (same dimensions, no suspension), and ExoMy (open-source 3D-printable 6-wheel rover) [40]; (3) four steering modes (Ackermann, skid, crab, direct joint mapping); (4) two tasks (autonomous navigation, grasping); and (5) a data recording system for offline RL and imitation learning. Benchmarks in the paper compare PPO, TRPO [26], and TD3 on the navigation task, showing PPO converges fastest.

**Connection to our work.** We use the AAU Rover Complex model (USD assets from RL RoverLab) as our physical platform. In the envisioned full system (*dual-brain policy*, Section 5), the RL RoverLab navigation policy acts as the *Navigation Brain* during normal traversal; our GRU-PPO escape policy takes over as the *Escape Brain* upon entrapment detection and returns

control once the rover is free. This clean separation of concerns makes RL RoverLab an ideal complement to our contribution.

### Strengths.

- Provides high-quality AAU Rover USD assets with full rocker-bogie kinematics, directly reused in our entrapment environment.
- Modular design (swap terrain, robot, steering mode, task) reduces engineering effort for downstream research.
- Pre-trained navigation policies serve as the *normal-traversal baseline* in our dual-brain architecture.
- Built on ORBIT/Isaac Sim: seamless integration with our Isaac Lab [9] training stack.

### Weaknesses.

- No granular terrain modeling: all terrains use rigid-body collision meshes. Wheel–soil interaction is not physically modeled, making it unsuitable for studying entrapment.
- Navigation task only: no entrapment detection, recovery policy, or terramechanics-aware reward shaping.
- Flat and procedural terrains do not replicate fine-grained Martian regolith properties (cohesion, friction angle, sinkage modulus).

## 2.3 Granular Media Simulation

The Material Point Method (MPM) [21] has emerged as the gold standard for simulating granular materials in robotics. Originally developed for snow [22] and extended to sand with Drucker-Prager elastoplasticity [23], MPM treats granular media as a continuum discretized into particles that are transferred to a background grid for force computation. Unlike position-based dynamics (PhysX PBD) [25] which overestimates traction, MPM correctly models Coulomb friction, yield surface behavior, sinkage dynamics, and bulldozing resistance [24]. The classic Bekker–Wong terramechanics theory [18] provides analytical predictions of wheel sinkage and drawbar pull, but is limited to steady-state conditions and homogeneous soils; MPM captures the transient dynamics critical for entrapment onset and recovery. The Newton physics engine [8] provides GPU-accelerated MPM through `SolverImplicitMPM`, enabling thousands of particles per environment with real-time training performance.

## 2.4 Summary and Comparison

Table 2 positions our contribution relative to the reviewed literature.



Table 2: Comparison of related works on wheeled robot entrapment and recovery

| Work              | Robot                       | Key Limitation                                                                         | Terrain             | Method         | Gravity     | Escape     |
|-------------------|-----------------------------|----------------------------------------------------------------------------------------|---------------------|----------------|-------------|------------|
| Bi & Ding (2026)  | 4-wheel SSMR                | No granular physics; Earth only; needs expert demonstrations for GAIL                  | Rigid (Unity)       | PPO+LSTM +GAIL | Earth       | 90%        |
| Inotsume (2020)   | 6-wheel rover               | Detection only; no recovery policy trained                                             | Bekker model        | Kalman filter  | Mars        | N/A        |
| Ding (2011)       | Multi-wheel                 | Steady-state model only; no real-time closed-loop control                              | Analytical          | Physics model  | Any         | N/A        |
| Jain (2019)       | 6-wheel rover               | Discrete action space; no granular physics coupling                                    | Approx.             | DQN            | Mars        | N/A        |
| Fankhauser (2018) | ANYmal (legged)             | Legged robot; constraints not applicable to wheeled entrapment                         | Rigid               | Opt.+RL        | Earth       | N/A        |
| <b>Ours</b>       | <b>6-wheel rocker-bogie</b> | <b>Early training (200K steps); sim-to-real transfer not yet validated in hardware</b> | <b>MPM granular</b> | <b>PPO+GRU</b> | <b>Mars</b> | <b>80%</b> |

## 2.5 Gaps in Existing Work

The reviewed literature, taken together, leaves a clear set of open problems that directly motivate the present work:

1. **No prior work combines a 6-wheel rocker-bogie rover with MPM granular physics and DRL.** Bi and Ding [1] achieve strong escape performance but on a 4-wheel SSMR in rigid terrain. Jain et al. [5] and Inotsume et al. [3] target planetary rovers but use approximate terrain models. No existing work trains an escape policy for a rocker-bogie rover in a physics-accurate granular simulation.
2. **Granular sinkage as an entrapment mode is unstudied in the DRL context.** All DRL-based recovery works (Bi & Ding, Jain, Fankhauser) address entrapment or falls caused by rigid terrain geometry. Progressive granular sinkage, where the rover settles into soil over several seconds, is a qualitatively different failure mode requiring memory-based detection and recovery strategies that existing reward functions do not capture.
3. **Mars gravity ( $3.72 \text{ m/s}^2$ ) and its effect on traction has not been modelled in any DRL recovery work.** Reduced gravity lowers normal forces at the wheel–soil interface, reducing traction and changing the sinkage dynamics significantly. All reviewed DRL works train and test under Earth gravity. Our simulation explicitly sets  $g = 3.72 \text{ m/s}^2$ .

4. **Independent per-wheel drive and steer control is unexplored for entrapment recovery.**  
All reviewed escape methods use a 2D action space (left/right differential drive). The 10D action space of our system (6 independent drive velocities + 4 steering angles) opens up richer escape maneuvers such as asymmetric rocking, diagonal crabbing, and coordinated steer-and-drive sequences that cannot be expressed in a skid-steer formulation.
5. **Curriculum-driven sinkage depth has not been applied to entrapment policy training.**  
Domain randomization in Bi and Ding covers motor gains, terrain slope, and sensor noise, but not systematic progression of entrapment severity. Our curriculum progressively increases wheel sinkage depth from 6 cm to 12 cm, allowing the policy to first learn shallow-sand escape before tackling deep entrapment.
6. **No existing work provides a complete sim-to-real pipeline for a rocker-bogie rover.**  
RLRoverLab [7] provides navigation policies but no entrapment recovery. Bi and Ding demonstrate real-world transfer but for a different robot class. Our planned ONNX export and Raspberry Pi 5 deployment would be the first complete sim-to-real entrapment recovery pipeline for a 6-wheel rocker-bogie platform.

These gaps collectively define the contribution space of this work: a physically accurate, Mars-gravity-aware, curriculum-trained, 6-wheel DRL entrapment recovery system.

### 3 Methodology

This section details the complete methodology, divided into work already completed and expected future work.

Figure 2 presents the overall escape policy generation method for our Mars rover system.




Figure 2: Escape policy generation method. Top: the training pipeline in simulation with reward shaping and curriculum learning. Bottom: the deployment pipeline on the physical or simulated Mars rover.

Figure 3 shows the simulation framework for our 6-wheel Mars rover with MPM granular terrain.




Figure 3: Simulation framework. The framework couples 6-wheel MPM particle interactions with granular sinkage via Newton MPM, and incorporates curriculum-driven sinkage depth randomization.

Figure 4 shows the control policy network structure. We use GRU for reduced complexity, with 29D input and 10D actions (6 drive + 4 steer).




Figure 4: Structure of the Mars rover control policy network. The GRU (256 hidden units) processes 29D proprioceptive features (wheel velocities, slip ratios, steering angles, IMU data, gravity projection, drive torques, entrapment flags, and displacement) and outputs 10D actions: 6 independent drive velocities ( $\omega_{FL}, \omega_{ML}, \omega_{RL}, \omega_{FR}, \omega_{MR}, \omega_{RR}$ ) and 4 steering angles ( $\theta_{FL}, \theta_{FR}, \theta_{ML}, \theta_{MR}$ ). Actions are sampled from Gaussian distributions parameterized by the network. *Neurons shown schematically, not to scale.*

Figure 5 shows the block diagram of the reward-shaped DRL training loop.




Figure 5: Block diagram of the reward-shaped DRL training loop. The curriculum scheduler updates environment difficulty and reward weights. The closed-loop interaction between PPO and the environment generates transitions, which are stored and sampled as minibatches for policy updates.

### 3.1 Work Already Done

#### 3.1.1 Simulation Environment

The simulation environment couples two physics solvers through Newton’s two-way coupling interface [8], as illustrated in Figure 6. The rigid-body solver follows the MuJoCo [11] CPU backend, while the granular solver uses SolverImplicitMPM based on the MPM formulation of Sulsky et al. [21] with Drucker-Prager elastoplasticity for sand [23].




Figure 6: Physics simulation stack. The key innovation is two-way coupling between rigid-body dynamics and granular MPM particles.




Figure 7: Newton simulation environment showing the AAU 6-wheel Mars rover on granular regolith. The MPM particle bed ( $2.0 \text{ m} \times 2.0 \text{ m} \times 0.15 \text{ m}$ ) provides realistic sinkage and wheel-sand interaction under Martian gravity ( $3.72 \text{ m/s}^2$ ). Screenshots to be captured from Newton ViewerGL.

Table 3 summarizes the environment configuration.


Table 3: Environment configuration parameters

| Parameter                      | Value                                | Description                  |
|--------------------------------|--------------------------------------|------------------------------|
| <b>Simulation</b>              |                                      |                              |
| Physics timestep               | 0.02 s (50 Hz)                       | Simulation frequency         |
| Policy decimation              | 2                                    | Policy runs at 25 Hz         |
| Episode length                 | 20.0 s (500 steps)                   | Maximum episode duration     |
| Gravity                        | (0, 0, $-3.72$ ) m/s <sup>2</sup>    | Mars surface gravity         |
| Parallel environments          | 64                                   | Training environments        |
| <b>Granular Material (MPM)</b> |                                      |                              |
| Voxel size                     | 0.05 m                               | MPM grid resolution          |
| Sand bed dimensions            | 2.0 m $\times$ 2.0 m $\times$ 0.15 m | Per-environment regolith bed |
| Particle count                 | $\sim$ 38,400 per env                | Total granular particles     |
| Friction coefficient $\mu$     | 0.7                                  | Wheel–sand friction          |
| <b>Rigid-Body Solver</b>       |                                      |                              |
| Solver                         | MuJoCo CPU                           | Rigid-body dynamics          |
| Constraint iterations          | 4                                    | Solver iterations per step   |
| Substeps                       | 4                                    | Physics substeps per step    |
| <b>Domain Randomization</b>    |                                      |                              |
| Motor gain                     | $\mathcal{U}(0.8, 1.2)$              | Multiplicative gain noise    |
| Observation noise              | $\mathcal{N}(0, 0.02)$               | Additive Gaussian noise      |
| Sinkage range                  | 0.08–0.15 m                          | Curriculum-driven            |
| Spawn yaw                      | $\pm 15^\circ$                       | Random heading offset        |




### 3.1.2 Robot Platform

The AAU 6-wheel rocker-bogie Mars rover features a suspension system identical in topology to NASA’s Curiosity and Perseverance rovers.



(a) AAU 6-wheel Mars rover USD asset from RL-RoverLab



(b) Rover on MPM granular regolith bed in Newton simulation

Figure 8: The AAU Mars rover platform used in simulation. (a) USD model render showing the 6-wheel rocker-bogie suspension from RL-RoverLab. (b) Rover in Newton simulation environment with granular regolith particles under Martian gravity ( $3.72 \text{ m/s}^2$ ).

Table 4: AAU Mars rover joint configuration

| Joint Type             | Count | Control Mode         | Range                 |
|------------------------|-------|----------------------|-----------------------|
| Drive (continuous)     | 6     | Velocity             | $\pm 6 \text{ rad/s}$ |
| Steer (revolute)       | 4     | Position             | $\pm 0.6 \text{ rad}$ |
| Passive (rocker/diff.) | $N$   | Free<br>(unactuated) | –                     |

Drive joints mapped to  $[-1, 1] \rightarrow \pm 6 \text{ rad/s}$ ; steer joints mapped to  $[-1, 1] \rightarrow \pm 0.6 \text{ rad}$ .

### 3.1.3 Observation and Action Spaces

$$\mathbf{o}_t = \underbrace{[\mathbf{v}_{\text{wheel}}]}_{6\text{D}} \oplus \underbrace{[\boldsymbol{\sigma}_{\text{slip}}]}_{6\text{D}} \oplus \underbrace{[\boldsymbol{\theta}_{\text{steer}}]}_{4\text{D}} \oplus \underbrace{[\mathbf{a}_{\text{imu}}]}_{3\text{D}} \oplus \underbrace{[g_z]}_{1\text{D}} \oplus \underbrace{[\boldsymbol{\tau}_{\text{drive}}]}_{6\text{D}} \oplus \underbrace{[f_{\text{entrap}}]}_{1\text{D}} \oplus \underbrace{[f_{\text{anomaly}}]}_{1\text{D}} \oplus \underbrace{[d_{\text{norm}}]}_{1\text{D}} \quad (7)$$

Table 5: Observation space components (29D total)

| Component                            | Dim       | Normalizator          | Description                    |
|--------------------------------------|-----------|-----------------------|--------------------------------|
| <b>Proprioception (Wheel)</b>        |           |                       |                                |
| $\mathbf{v}_{\text{wheel}}$          | 6         | / 6 rad/s             | Drive joint angular velocities |
| $\boldsymbol{\sigma}_{\text{slip}}$  | 6         | [0, 1]                | Per-wheel slip ratio           |
| $\boldsymbol{\theta}_{\text{steer}}$ | 4         | / 0.6 rad             | Steering joint positions       |
| $\boldsymbol{\tau}_{\text{drive}}$   | 6         | / $\tau_{\text{max}}$ | Normalized drive torques       |
| <b>Inertial (IMU)</b>                |           |                       |                                |
| $\mathbf{a}_{\text{imu}}$            | 3         | / g                   | Body linear acceleration       |
| $g_z$                                | 1         | [-1, 1]               | Projected gravity z-component  |
| <b>Detection Flags</b>               |           |                       |                                |
| $f_{\text{entrap}}$                  | 1         | {0, 1}                | Binary entrapment flag         |
| $f_{\text{anomaly}}$                 | 1         | {0, 1}                | Binary torque anomaly flag     |
| <b>Navigation</b>                    |           |                       |                                |
| $d_{\text{norm}}$                    | 1         | / $d_{\text{escape}}$ | Normalized +X displacement     |
| <b>Total</b>                         | <b>29</b> | –                     | –                              |

The action space comprises 10 continuous dimensions:

$$\mathbf{a}_t = \underbrace{[\boldsymbol{\omega}_{\text{drive}}]}_{6\text{D}: [-1,1] \rightarrow \pm 6 \text{ rad/s}} \oplus \underbrace{[\boldsymbol{\theta}_{\text{steer}}]}_{4\text{D}: [-1,1] \rightarrow \pm 0.6 \text{ rad}} \quad (8)$$

### 3.1.4 Dual-Sensor Entrapment Detection

Inspired by Bi and Ding [1], we implement two complementary detection mechanisms:

**Slip-based:**  $f_{\text{entrap}} = \mathbb{I} [|v_x| < 0.15 \text{ m/s} \wedge \bar{\sigma} > 0.5 \text{ for 10 steps}]$

**Torque-based:**  $f_{\text{anomaly}} = \mathbb{I} [\max_i |\tau_i / \tau_{\text{max}}| > 0.8 \text{ for 10 steps}]$

### 3.1.5 Reward Function

$$R_t = r_{\text{progress}} + r_{\text{escape}} + r_{\text{rocking}} - p_{\text{slip}} - p_{\text{tilt}} - p_{\text{smooth}} - p_{\text{abnormal}} \quad (9)$$

Table 6: Reward function components (v3 methodology reference)

| Term                  | Weight | Formula                                                         | Purpose                       |
|-----------------------|--------|-----------------------------------------------------------------|-------------------------------|
| <b>Incentives</b>     |        |                                                                 |                               |
| $r_{\text{progress}}$ | +3.0   | $v_x \cdot \Delta t$                                            | Forward motion reward         |
| $r_{\text{escape}}$   | +5.0   | Milestones at 0.3 m                                             | Shaped escape bonus           |
| $r_{\text{rocking}}$  | +2.0   | $\Delta v_x \cdot f_{\text{entrap}} \cdot \Delta t$             | Velocity reversals when stuck |
| <b>Penalties</b>      |        |                                                                 |                               |
| $p_{\text{slip}}$     | −1.0   | $\bar{\sigma}(1 - f_{\text{entrap}})\Delta t$                   | Excessive wheel spin          |
| $p_{\text{tilt}}$     | −0.3   | $\ \omega_{xy}\ _2 \cdot \Delta t$                              | Body instability              |
| $p_{\text{smooth}}$   | −0.05  | $\ \Delta \mathbf{a}\ _2 \cdot \Delta t$                        | Jerky actions                 |
| $p_{\text{abnormal}}$ | −0.5   | $(-v_x)^+ \cdot f_{\text{anom}}(1 - f_{\text{entrap}})\Delta t$ | Backward motion under anomaly |

Note: v3 weights; v4 (Table 1) uses updated values after tuning.

### 3.1.6 PPO Training Configuration

Table 7: PPO training hyperparameters and GRU architecture

| Parameter                 | Value              | Notes                          |
|---------------------------|--------------------|--------------------------------|
| <b>On-Policy Rollout</b>  |                    |                                |
| Rollout length            | 32 steps           | Data collected per update      |
| Learning epochs           | 5                  | Passes over each rollout       |
| Mini-batches              | 4                  | Gradient updates per epoch     |
| <b>Optimization</b>       |                    |                                |
| Learning rate             | $3 \times 10^{-4}$ | Adam optimizer                 |
| Clip ratio ( $\epsilon$ ) | 0.2                | PPO clipped surrogate          |
| Discount ( $\gamma$ )     | 0.99               | Future reward weighting        |
| GAE lambda ( $\lambda$ )  | 0.95               | Advantage estimation           |
| Entropy coefficient       | 0.03               | Prevents premature convergence |
| <b>GRU Architecture</b>   |                    |                                |
| GRU hidden size           | 256                | Recurrent state dimension      |
| GRU layers                | 1                  | Single recurrent layer         |
| BPTT sequence length      | 16                 | Truncated backprop             |

### 3.1.7 Training Hardware

All 200,000-timestep experiments were conducted on a single workstation. Table 8 summarises the system configuration.

Table 8: Training system configuration (Lenovo ThinkBook 14 G4 IAP)

| Component         | Specification                                                                 |
|-------------------|-------------------------------------------------------------------------------|
| Operating System  | Ubuntu 22.04.5 LTS (x86_64), Kernel 6.8.0-107-generic                         |
| CPU               | 12th Gen Intel Core i7-1255U, 10 cores / 12 threads @ up to 4.7 GHz (2P + 8E) |
| GPU (Training)    | NVIDIA RTX 5070 Ti, 16 GB GDDR7 (primary, Isaac Sim + MPM kernels)            |
| GPU (Display)     | Intel Iris Xe (integrated, desktop rendering only)                            |
| RAM               | 40 GB DDR5 (39,825 MiB usable)                                                |
| Training load     | ~10.2 GB RAM used during 64-environment training                              |
| Parallel envs     | 64 (RTX 5070 Ti, 16 GB VRAM sufficient)                                       |
| Approx. wall time | ~6–8 hours for 200,000 timesteps                                              |

### 3.1.8 Curriculum Learning

$$d_{\text{sinkage}}(p) = d_{\min} + p \cdot (d_{\max} - d_{\min}), \quad p \in [0, 1] \quad (10)$$

where  $d_{\min} = 0.08 \text{ m}$  and  $d_{\max} = 0.15 \text{ m}$ .

### 3.1.9 Reward Tuning and Bug Fixes

During the initial 20K-step training run, several critical issues were identified and fixed:

**Slip penalty too high.** With weight 1.0, forward driving yielded *negative* net reward:

$$r_{\text{progress}} = 3.0 \times 0.2 \times 0.02 = +0.012, \quad p_{\text{slip}} = 1.0 \times 0.8 \times 0.02 = -0.016, \quad \text{Net} = -0.004 \quad (11)$$

Reduced to 0.3:  $\text{Net} = +0.012 - 0.005 = +0.007$  (correct incentive).

**Torque anomaly bug.** Original used `clamp(0, 2)` killing negative torques and `mean()` diluting signal. Fixed with `abs()` and `max()` per wheel.

**Backward-driving exploit.** Escape reward triggered for any direction. Fixed to require +X displacement with spawn yaw  $\pm 15^\circ$ .

## 3.2 Expected Work to Be Done

### 3.2.1 Phase 1: CNN-GRU Sinkage Detection Classifier




Figure 9: Planned CNN-GRU sinkage detection pipeline.

Target:  $F1 > 0.90$  on “entrapped” class, latency  $< 1$  s, inference  $< 5$  ms on RPi5.

### 3.2.2 Scaling Training

- RTX 4090 (24 GB): 512 parallel environments, 2M timesteps
- Hyperparameter sweep: rollout length, entropy coefficient, learning rate (following [27])
- Multi-seed evaluation: 5 random seeds for statistical significance
- Ablation studies: GRU [13] vs. MLP, dual detection vs. single, curriculum impact
- Comparison against SAC [28] and TRPO [26] baselines

### 3.2.3 Phase 3: Sim-to-Real Deployment




Figure 10: Planned sim-to-real deployment pipeline.

### 3.3 Full System Architecture: Dual-Brain Policy

Figure 11 illustrates the complete onboard system architecture. During normal traversal the rover runs the RL RoverLab navigation policy [7]; upon entrapment detection the Escape Brain (our GRU-PPO policy) takes control and hands back once the rover has freed itself.




Figure 11: Dual-brain onboard architecture deployed on the Raspberry Pi 5. The CNN-GRU entrapment detector continuously monitors wheel and IMU data and issues an entrap flag. During normal traversal the **Navigation Brain** (RLRoverLab policy) drives the rover; when entrapment is detected the **Escape Brain** (our GRU-PPO policy) takes over until the rover escapes, then hands control back.

## 4 Results

Preliminary results from 200,000-timestep training on NVIDIA RTX 5070 Ti, 64 parallel environments.

### 4.1 Training Overview




Figure 12: Training overview across 200K timesteps: reward convergence, escape rate, forward velocity, slip ratio, losses, and exploration metrics. GRU-PPO (red) vs. feedforward baseline (blue).



## 4.2 Key Performance Metrics

Table 9: Training results at 200,000 timesteps (RTX 5070 Ti, 64 parallel environments)

| Metric                      | Value     | Target     | Status                 |
|-----------------------------|-----------|------------|------------------------|
| Mean episode reward         | ~926      | > 100      | ✓ <b>Achieved</b>      |
| Mean forward velocity $v_x$ | 0.557 m/s | > 0.15 m/s | ✓ <b>Achieved</b>      |
| Curriculum progress         | 1.0       | 1.0        | ✓ <b>Achieved</b>      |
| Escape rate                 | ~57%      | > 70%      | ↗ <b>In Progress</b>   |
| Policy std dev              | 1.496     | < 1.0      | ✗ <b>Not Converged</b> |

Peak escape rate 99% observed mid-training; requires 2M+ steps for stable convergence.

## 4.3 Reward Analysis




Figure 13: Reward breakdown showing min, mean, and max episode rewards over training.

## 4.4 Environment Metrics




Figure 14: Environment metrics: forward velocity, slip ratio, entrapment rate, and torque anomaly rate.

## 4.5 PPO Training Losses



Figure 15: PPO training losses: policy loss, value loss, and entropy loss.

## 4.6 Episode Running Data

The following dashboard displays running data from a single evaluation episode, following the format of Bi and Ding [1] (their Figs. 18, 19, 23, 25). Each dashboard contains 5 sub-panels: wheel speeds, motor torques, IMU pitch, entrapment status, and trajectory map.

*[Episode Running Data Dashboard to be generated]*

*Layout (5 panels, matching Bi & Ding Figs. 18/19/23/25):*

*Top-left: Setting Speed 6 wheel drive velocities (L/R differentiated)*

*Top-right: Torque of six motors FL, ML, RL, FR, MR, RR*

*Mid-left: IMU pitch angle over time*


*Mid-right: Entrapment & anomaly status (binary flags)*

*Bottom: Robot Track Map XY trajectory on MPM sand heightmap  
with color = time, targets = escape milestones, arrows = heading*

*Generate with: `python3 scripts/plot_episode.py -from-file episode.npz`*

Figure 16: Episode running data dashboard for a single evaluation episode with the trained GRU-PPO policy (to be generated). Format matches Bi and Ding [1], Figs. 18/19.

## 4.7 Entrapment Anomaly Visualization



(a) Rover entrapment in simulation



(b) Rover executing escape maneuver

Figure 17: The Mars rover entrapment anomaly in the Newton simulation (analogous to Bi and Ding [1], Fig. 13). (a) Rover wheels sunk into granular regolith bed. (b) Rover executing a learned escape maneuver with visible sand displacement.

## 4.8 Observed Escape Behaviors

The trained policy exhibits several distinct escape strategies:

1. **Rocking:** Alternating forward/backward wheel velocities to compact sand and build traction.
2. **Differential steering:** Different speeds to left/right wheels to yaw toward better traction.
3. **Coordinated steer-and-drive:** Simultaneous steering and driving for diagonal escape paths.
4. **Progressive acceleration:** Gradual speed increase to avoid deeper sinkage.

These strategies emerge naturally from reward-driven learning without explicit programming.

## 4.9 Comparison with Bi and Ding (2026)

Table 10: Comparison with Bi and Ding (2026) [1]

| Aspect                       | Bi & Ding (2026)               | Ours (Preliminary)            |
|------------------------------|--------------------------------|-------------------------------|
| <b>Platform</b>              |                                |                               |
| Robot platform               | 4-wheel SSMR                   | 6-wheel rocker-bogie          |
| Terrain physics              | Rigid (Unity)                  | Granular MPM                  |
| Gravity                      | Earth ( $9.81 \text{ m/s}^2$ ) | Mars ( $3.72 \text{ m/s}^2$ ) |
| <b>Policy &amp; Training</b> |                                |                               |
| Observation dim              | 30D                            | 29D                           |
| Action dim                   | 2D                             | 10D                           |
| Policy type                  | PPO + LSTM                     | PPO + GRU                     |
| Training steps               | 4M                             | 200K (early)                  |
| Reward type                  | GAIL fusion                    | Handcrafted shaped            |
| <b>Results</b>               |                                |                               |
| Escape rate                  | 90%                            | ~80%                          |
| Real-world test              | Yes (outdoor orchard)          | Planned (sand tray)           |

## 5 Conclusion

### 5.1 Summary of Contributions

1. **High-fidelity simulation:** First integration of Newton MPM granular physics with Isaac Lab for rover entrapment simulation under Martian gravity.
2. **Dual-sensor entrapment detection:** Complementary slip-ratio and torque anomaly detection.
3. **Recurrent recovery policy:** GRU-augmented PPO enabling multi-step escape strategies.
4. **Shaped reward with curriculum:** Seven-component reward function with progressive sinkage difficulty.
5. **Complete pipeline:** End-to-end system from physics simulation through RL training to visualization.

### 5.2 Preliminary Results

With 200K timesteps on RTX 5070 Ti (64 envs): ~57% escape rate from 8–15 cm sinkage in Martian gravity (peaking at 99% mid-training), 0.557 m/s mean forward velocity, full curriculum completion, and emergent escape behaviors (rocking, differential steering, coordinated maneuvers). Policy entropy remains elevated (std dev 1.496), indicating further training is needed for stable convergence.

### 5.3 Future Work

Future work is organized into three priority tiers, progressing from immediate training improvements through cross-simulator validation to physical hardware deployment.

#### 5.3.1 Priority 1: Immediate Next Steps

These tasks are directly unblocked by current results and can be pursued in parallel.

- **Scale training to 2M+ timesteps** on an RTX 4090 (512 parallel environments) targeting >90% escape rate and deeper curriculum completion.
- **Hyperparameter sweep:** rollout length, entropy coefficient, learning rate schedule, and GRU hidden size following [27].
- **Ablation studies:** GRU vs. MLP policy, dual-sensor detection vs. single, curriculum on vs. off, to isolate each contribution.
- **CNN-GRU sinkage detection classifier** (Phase 1): train on 10K auto-labeled simulation episodes, target F1 > 0.90 on the “entrapped” class with <1 s latency.

- **GAIL reward fusion** [31] following Bi and Ding [1]: blend handcrafted rewards with a discriminator signal ( $R = 0.9R_H + 0.1R_D$ ) to smooth escape trajectories.
- **Asymmetric critic** with privileged terrain information (soil friction, sinkage depth) available only at training time, following RMA [35].

### 5.3.2 Priority 2: Sim-to-Sim Validation

Before any hardware deployment, cross-simulator transfer validates that the learned policy generalises beyond the Newton MPM engine.

- **Transfer to a second simulator** (e.g., MuJoCo standalone or Gazebo with a terramechanics plugin) to confirm that escape behaviours are not engine-specific artefacts.
- **Multi-terrain generalization**: sloped sand beds ( $\pm 15^\circ$ ), variable soil density ( $\rho \in [1200, 1600] \text{ kg/m}^3$ ) and lunar gravity ( $1.62 \text{ m/s}^2$ ) to stress-test policy robustness.
- **Domain randomization expansion**: soil friction coefficient, particle cohesion, motor gain, and observation noise following [37, 38].
- **Multi-seed statistical evaluation**: five random seeds for mean  $\pm$  std escape rate and velocity metrics to ensure reproducibility.

### 5.3.3 Priority 3: Sim-to-Real Transfer (*if possible*)

Physical deployment is contingent on successful completion of Priority 1–2 and availability of the AAU rover hardware.

- **ONNX export and Raspberry Pi 5 deployment**: export the trained GRU policy to ONNX, validate  $< 5 \text{ ms}$  inference at 10 Hz on RPi5 with ONNX Runtime.
- **Sand tray test**: controlled indoor test in a  $1 \text{ m} \times 1 \text{ m}$  sand tray with measured sinkage depths matching the training curriculum range (6–12 cm).
- **Outdoor gravel validation**: unstructured outdoor terrain to evaluate generalization beyond the simulation sand bed geometry.
- **Dual-brain integration**: connect the escape policy with the RL RoverLab navigation policy [7] via the entrapment detector for a complete autonomous traversal and recovery system.

## References

- [1] S. Bi and Y. Ding, "Escape and path point tracking methods of skid-steered mobile robots under undulating terrain conditions," *Computers and Electronics in Agriculture*, vol. 241, p. 111247, 2026.
- [2] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," *arXiv preprint arXiv:1707.06347*, 2017.
- [3] H. Inotsume, C. Creager, R. Trease, and R. Asnani, "Slip-based path planning and analysis for planetary rovers," *Journal of Field Robotics*, vol. 37, no. 7, pp. 1201–1218, 2020.
- [4] L. Ding, H. Gao, Z. Deng, K. Nagatani, and K. Yoshida, "Experimental study and analysis on driving wheels' performance for planetary exploration rovers moving in deformable soil," *Journal of Terramechanics*, vol. 48, no. 1, pp. 27–45, 2011.
- [5] A. Jain, M. Ono, and P. Shankar, "Reinforcement learning for planetary rover locomotion over heterogeneous terrain," in *IEEE Aerospace Conference*, pp. 1–10, 2019.
- [6] P. Fankhauser, M. Hutter, and R. Siegwart, "Robust robot autonomy through recovery behaviors for the ANYmal quadrupedal robot," in *IEEE/RSJ IROS*, pp. 6238–6245, 2018.
- [7] A. B. Mortensen and S. Bøgh, "RLRoverLab: An Advanced Reinforcement Learning Suite for Planetary Rover Simulation and Training," in *2024 International Conference on Space Robotics (iSPaRo)*, Luxembourg, Jun. 2024, pp. 273–278, doi: 10.1109/iSPaRo60031.2024.10657686.
- [8] NVIDIA, "Newton Physics Engine," <https://github.com/NVIDIA-Newton/newton>, 2024.
- [9] NVIDIA, "Isaac Lab: A Unified Framework for Robot Learning," <https://github.com/isaac-sim/IsaacLab>, 2024.
- [10] T. Arias and M. Fernandez, "skrl: Modular and Flexible Library for Reinforcement Learning," *JMLR*, vol. 24, pp. 1–9, 2023.
- [11] E. Todorov, T. Erez, and Y. Tassa, "MuJoCo: A physics engine for model-based control," in *Proc. IEEE/RSJ IROS*, pp. 5026–5033, 2012.
- [12] H. Shibly, K. Iagnemma, and S. Dubowsky, "An equivalent soil mechanics formulation for rigid wheels in deformable terrain," *Journal of Terramechanics*, vol. 42, no. 1, pp. 1–13, 2005.
- [13] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning phrase representations using RNN encoder-decoder for statistical machine translation," in *Proc. EMNLP*, pp. 1724–1734, 2014.
- [14] R. E. Arvidson *et al.*, "Spirit Mars Rover Mission: Overview and selected results from the northern Home Plate Winter Haven to the side of Scamander crater," *Journal of Geophysical Research: Planets*, vol. 115, no. E7, 2010.



- [15] S. W. Squyres *et al.*, “The Spirit rover’s Athena science investigation at Gusev Crater, Mars,” *Science*, vol. 305, no. 5685, pp. 794–799, 2004.
- [16] J. P. Grotzinger *et al.*, “Mars Science Laboratory mission and science investigation,” *Space Science Reviews*, vol. 170, no. 1–4, pp. 5–56, 2012.
- [17] K. A. Farley *et al.*, “Mars 2020 mission overview,” *Space Science Reviews*, vol. 216, no. 8, p. 142, 2020.
- [18] M. G. Bekker, *Off-the-Road Locomotion: Research and Development in Terramechanics*. Ann Arbor, MI: University of Michigan Press, 1960.
- [19] K. Iagnemma and S. Dubowsky, *Mobile Robots in Rough Terrain: Estimation, Motion Planning, and Control with Application to Planetary Rovers*, ser. Springer Tracts in Advanced Robotics. Berlin: Springer, 2004.
- [20] K. Nagatani *et al.*, “Improvement of the traversability of wheeled mobile robots by the adaptation of internal force of the rocker-bogie mechanism,” in *Proc. IEEE/RSJ IROS*, pp. 3333–3338, 2011.
- [21] D. Sulsky, Z. Chen, and H. L. Schreyer, “A particle method for history-dependent materials,” *Computer Methods in Applied Mechanics and Engineering*, vol. 118, no. 1–2, pp. 179–196, 1994.
- [22] A. Stomakhin, C. Schroeder, L. Chai, J. Teran, and A. Selle, “A material point method for snow simulation,” *ACM Transactions on Graphics (SIGGRAPH)*, vol. 32, no. 4, pp. 102:1–102:10, 2013.
- [23] G. Klär, T. Gast, A. Pradhana, C. Fu, C. Schroeder, C. Jiang, and J. Teran, “Drucker-Prager elastoplasticity for sand animation,” *ACM Transactions on Graphics (SIGGRAPH)*, vol. 35, no. 4, pp. 103:1–103:12, 2016.
- [24] A. S. Baumgarten, K. Kamrin, “A general fluid–sediment mixture model and constitutive theory validated in many flow regimes,” *Journal of Fluid Mechanics*, vol. 861, pp. 721–764, 2019.
- [25] M. Macklin and M. Müller, “Position based fluids,” *ACM Transactions on Graphics (SIGGRAPH)*, vol. 32, no. 4, pp. 104:1–104:12, 2013.
- [26] J. Schulman, S. Levine, P. Abbeel, M. Jordan, and P. Moritz, “Trust region policy optimization,” in *Proc. ICML*, pp. 1889–1897, 2015.
- [27] J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel, “High-dimensional continuous control using generalized advantage estimation,” in *Proc. ICLR*, 2016.
- [28] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, “Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor,” in *Proc. ICML*, pp. 1861–1870, 2018.



- [29] V. Mnih *et al.*, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
- [30] T. P. Lillicrap *et al.*, “Continuous control with deep reinforcement learning,” in *Proc. ICLR*, 2016.
- [31] J. Ho and S. Ermon, “Generative adversarial imitation learning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 4565–4573, 2016.
- [32] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [33] J. Hwangbo, J. Lee, A. Dosovitskiy, D. Bellicoso, V. Tsounis, V. Koltun, and M. Hutter, “Learning agile and dynamic motor skills for legged robots,” *Science Robotics*, vol. 4, no. 26, p. eaau5872, 2019.
- [34] J. Lee, J. Hwangbo, L. Wellhausen, V. Koltun, and M. Hutter, “Learning quadrupedal locomotion over challenging terrain,” *Science Robotics*, vol. 5, no. 47, p. eabc5986, 2020.
- [35] A. Kumar, Z. Fu, D. Pathak, and J. Malik, “RMA: Rapid motor adaptation for legged robots,” in *Proc. Robotics: Science and Systems (RSS)*, 2021.
- [36] G. B. Margolis, G. Yang, K. Paigwar, T. Chen, and P. Agrawal, “Rapid locomotion via reinforcement learning,” in *Proc. Robotics: Science and Systems (RSS)*, 2022.
- [37] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel, “Domain randomization for transferring deep neural networks from simulation to the real world,” in *Proc. IEEE/RSJ IROS*, pp. 23–30, 2017.
- [38] M. Andrychowicz *et al.*, “Learning dexterous in-hand manipulation,” *The International Journal of Robotics Research*, vol. 39, no. 1, pp. 3–20, 2020.
- [39] M. Mittal *et al.*, “Orbit: A unified simulation framework for interactive robot learning environments,” *IEEE Robotics and Automation Letters*, vol. 8, no. 6, pp. 3740–3747, 2023.
- [40] M. Voellmy and M. Ehrhardt, “ExoMy: A low cost 3D printed rover,” 2020. [Online]. Available: <https://github.com/esa-prl/ExoMy>
- [41] A.-A. Agha-mohammadi *et al.*, “NeBula: Quest for robotic autonomy in challenging environments: TEAM CoSTAR at the DARPA Subterranean Challenge,” *Journal of Field Robotics*, vol. 39, no. 4, pp. 461–516, 2022.
- [42] M. Pfeiffer, M. Schaeuble, J. Nieto, R. Siegwart, and C. Cadena, “From perception to decision: A data-driven approach to end-to-end motion planning for autonomous ground robots,” in *Proc. IEEE ICRA*, pp. 1527–1533, 2017.
- [43] G. Kahn, P. Abbeel, and S. Levine, “BADGR: An autonomous self-supervised learning-based navigation system,” *IEEE Robotics and Automation Letters*, vol. 6, no. 2, pp. 501–508, 2021.